HEXING

Intelligence and Innovation Driving Results

AI Server Maintenance Manual

Version 1.0.0

Qiao Lei

August 19, 2025

Contents

Abstract

This article delivers comprehensive user guides and best-practice maintenance recommendations for the array of AI-enabled services hosted on the test center's AI server. It is crafted for system integrators and operations staff, equipping them to efficiently support these systems and to extend the platform with advanced applications.

Keywords:

AI, LLM, Open WebUI, ollama, Ragflow, Dify, Agent, OCR, PaddlePaddle, AI4Test, Common Language .....

1 Server Information

1.1 Hardware basic information

- **Type:** DELL Precision 3660 Workstation
- **Service code:** 35K0L24
- **CPU:** 13th Gen Intel(R) Core(TM) i9-13900K
- **Video Card:** NVIDIA GeForce RTX 4090 Memory: 24G
- **Memory:** 128G
- **Hard disk 1:** nvme0n1 1T (NVMe SSD)
- **Hard disk 2:** sda 4T

1.2 Login Address

- **Server address:** 172.16.33.244
- **Admin User name:** Administrator, **Password:** hexing@2025
- **User name** : ems, **Password:** sme8003

2 AI Related System

2.1 Ollama - AI local LLM inference server

How to **install** the ollama on Ubuntu ?

```
# download the installation package
curl -L <https://ollama.com/download/ollama-linux-amd64.tgz> -o ollama-
    linux-amd64.tgz

# uncompress the package
cd /home/user/tools/ollama
```

```
gunzip ollama-linux-amd64.tgz
tar xvf ollama-linux-amd64.tar

# run the ollama service, also can configure to the system startup
   service
cd ollama/bin
./ollama serve
```

By typing in the browser"**http://172.16.33.244:11434**" to check if Ollama's service is working properly

if display "**Ollama is running**" that means Ollama is running well.

After **"ollama pull"** command, can use **"ollama list"** to check the current models.

```
(.venv) ems@aisrv:~$ ollama -v
ollama version is 0.11.4
(.venv) ems@aisrv:~$ ollama list
NAME                                        ID              SIZE      MODIFIED
gpt-oss:20b                                 e95023cf3b7b    13 GB     2 days ago
qwen3:32b                                   e1c9f234c6eb    20 GB     3 months ago
qwen3:30b                                   2ee832bc15b5    18 GB     3 months ago
gemma3:27b                                  30ddded7fba6    17 GB     4 months ago
gemma3:12b                                  6fd036cefda5    8.1 GB    5 months ago
olmOCR-7B:latest                            2c70ee7003bc    15 GB     5 months ago
maryasov/qwen2.5-cline:latest               38aafeb8f691    12 GB     5 months ago
deepseek-coder:33b                          acec7c0b0fd9    18 GB     5 months ago
huihui_ai/deepseek-r1-abliterated:32b       fb53b3296912    19 GB     6 months ago
deepseek-r1:32b                             38056bbcbb2d    19 GB     6 months ago
phi4:14b                                    ac896e5b8b34    9.1 GB    7 months ago
qwen2.5-coder:32b                           4bd6cbf2d094    19 GB     9 months ago
llama3.2:3b                                 a80c4f17acd5    2.0 GB    10 months ago
qwen2.5:32b                                 9f13ba1299af    19 GB     10 months ago
codegemma:code                              926331004170    1.6 GB    11 months ago
yi-coder:9b                                 0eed9e7baf59    5.0 GB    11 months ago
mistral-nemo:12b                            4b300b8c6a97    7.1 GB    12 months ago
t1c/deepseek-math-7b-rl:latest              5a7b7fcff0c2    4.2 GB    12 months ago
llama3.1:latest                             62757c860e01    4.7 GB    12 months ago
mxbai-embed-large:latest                    468836162de7    669 MB    13 months ago
nomic-embed-text:latest                     0a109f422b47    274 MB    13 months ago
deepseek-coder-v2:latest                    8577f96d693e    8.9 GB    13 months ago
gemma2:27b                                  53261bc9c192    15 GB     13 months ago
```

Figure 1: All the available LLMs

By "**nvidia-smi**" to check video card information

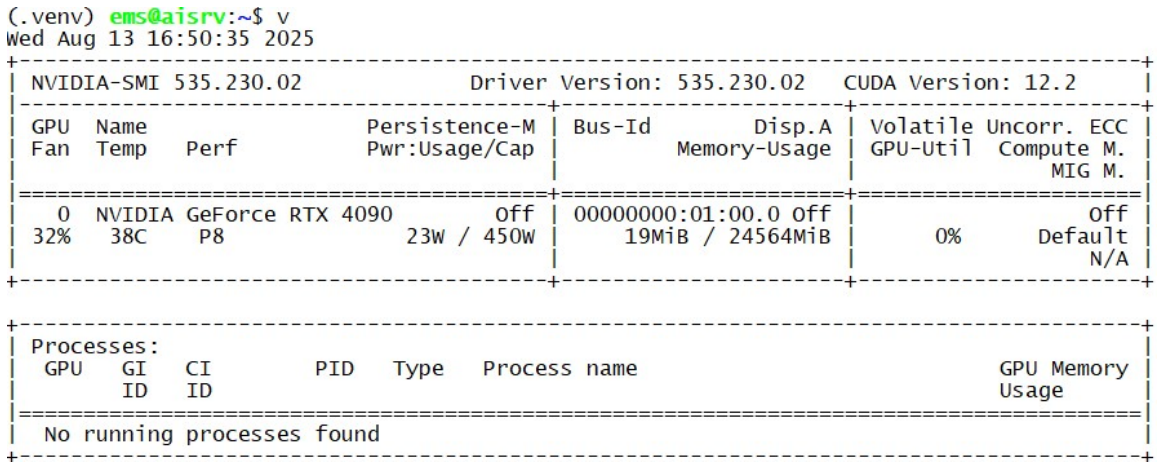

Figure 2: video card information

2.2 Open WebUI - web interface for LLM chat

Run **Open WebUI** via Docker, and be sure to use the **CUDA-enabled** version because it needs to run on the **GPU**.

Note:

1. The **CUDA** version is required.
2. The external default service port is **3000**, but it can be changed.
3. he container service starts and stops **automatically** in sync with the server's startup and shutdown.
4. If pulling the Docker image fails or is slow, switch to a different **image mirror** or use a **VPN**.

```
docker run -d -p 3000:8080 --gpus all --add-host=host.docker.internal:
   host-gateway -v open-webui:/app/backend/data --name open-webui --
   restart always ghcr.io/open-webui/open-webui:cuda
```

Login to Open Webui at **http://172.16.33.244:3000** . When you log in for the first time, you need to create a **administrator account** and password. Subsequent users can register themselves, and the system will **automatically approve** the registrations (you can optionally configure the approval to require a administrator's confirmation instead). Currently, the admin user name is : **q30china@gmail.com**, password: **hexing@2025**

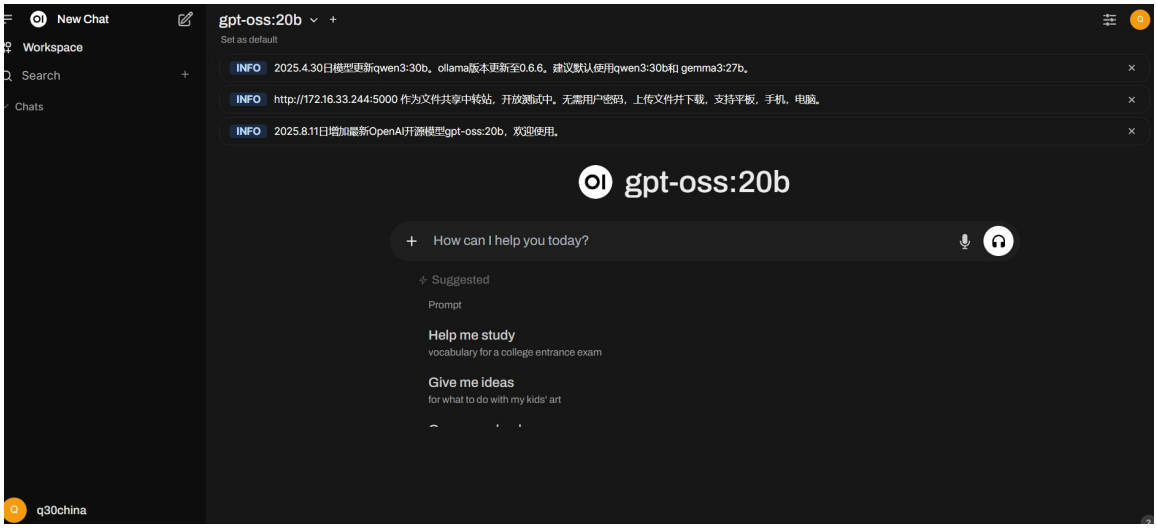

Figure 3: Open Webui User Interface

2.3 Ragflow - Retrieval-Augmented Generation (RAG) systems

Run all the container images with Docker. Once they start normally, the service will be exposed on the default port 8080 . During the first initialization you must set up a system-administrator account and password. Currently, the admin user name is : **q30china@gmail.com**, password: **hexing@2025**

```
# start the service
cd /data/git/ragflow/docker
docker compose -f docker-compose-gpu.yml up -d

# stop the service
docker compose -f docker-compose-gpu.yml down
```

Figure 4: Ragflow User Interface

How to use the Ragflow local knowledge database, and create a web application to use this.

step 1: Add local models service (Ollama is running)

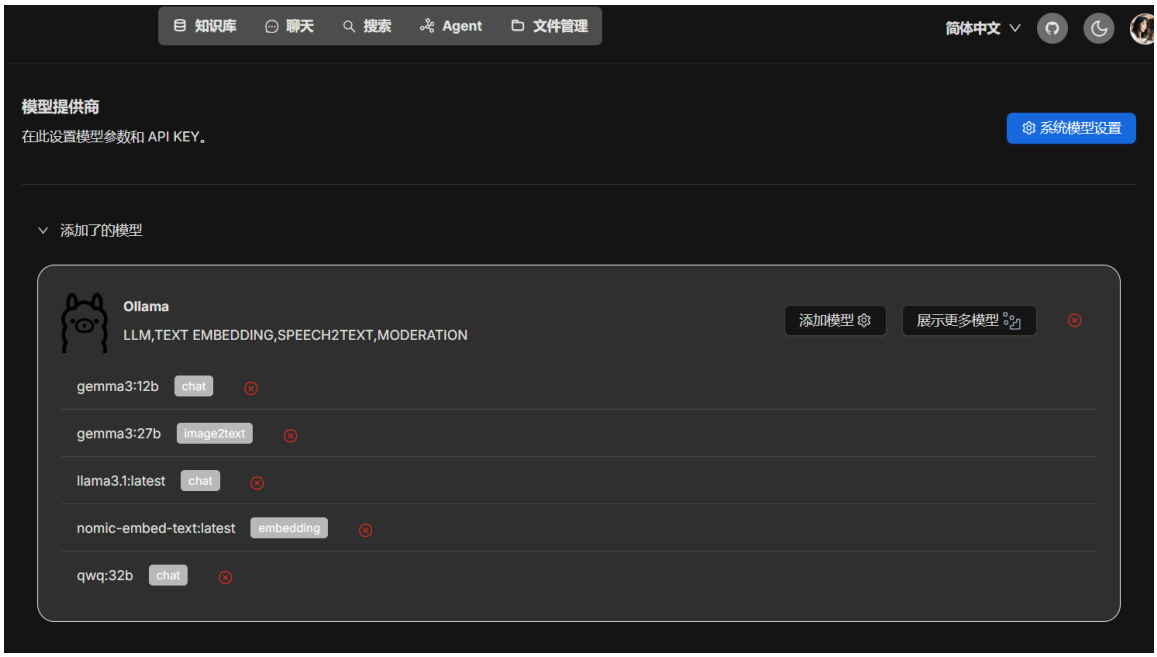

Figure 5: Add local models

step 2: Create a new knowledge base, upload local documents, and wait until the system has finished parsing and vectorizing the storage before proceeding.

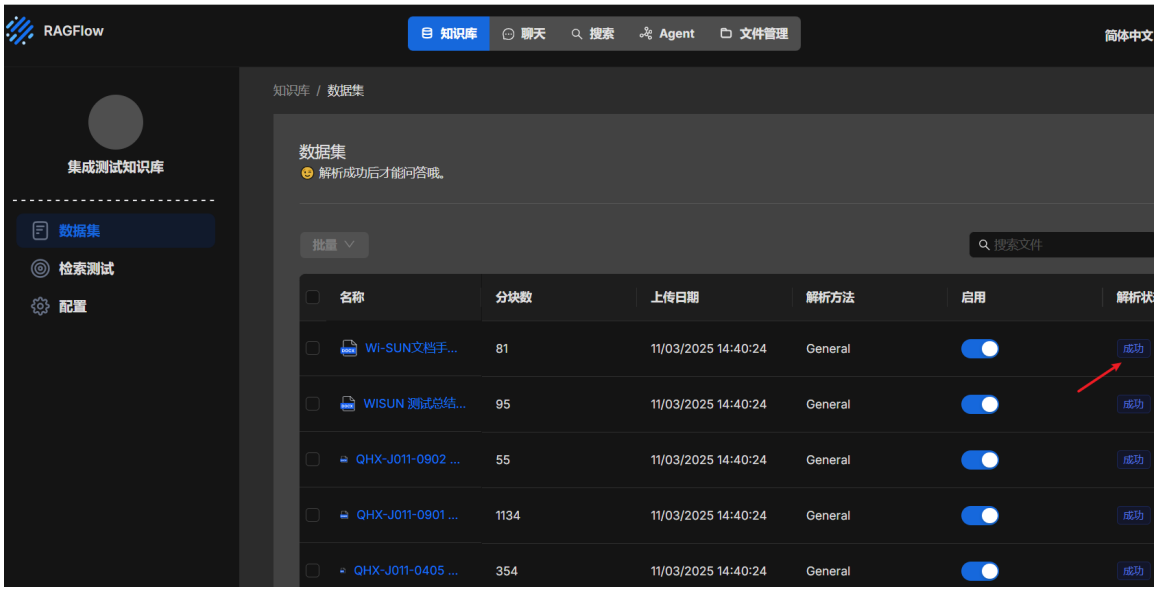

Figure 6: Upload local files to knowledge database

step 3: In the knowledge-base configuration screen, be sure to select the document's language and choose the embedded vector model (pulled via Ollama). Adjust any other parameters as needed, but typically the defaults are fine.

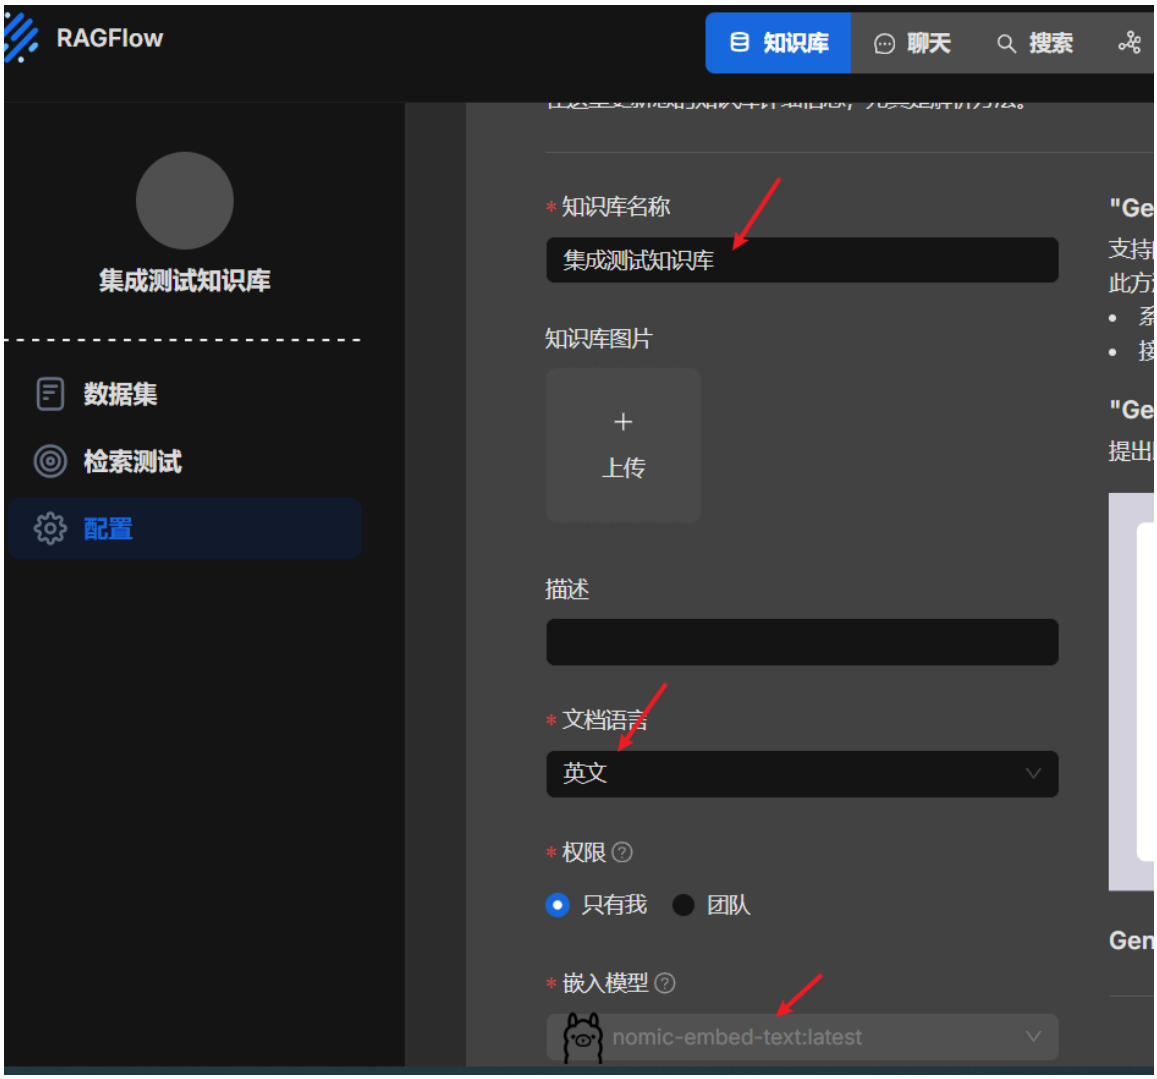

Figure 7: Configuration of knowlege database

step 4: Create a new digital-human application for new-employee training, using the locally created knowledge base and configuring the LLM to a model from the local Ollama service.

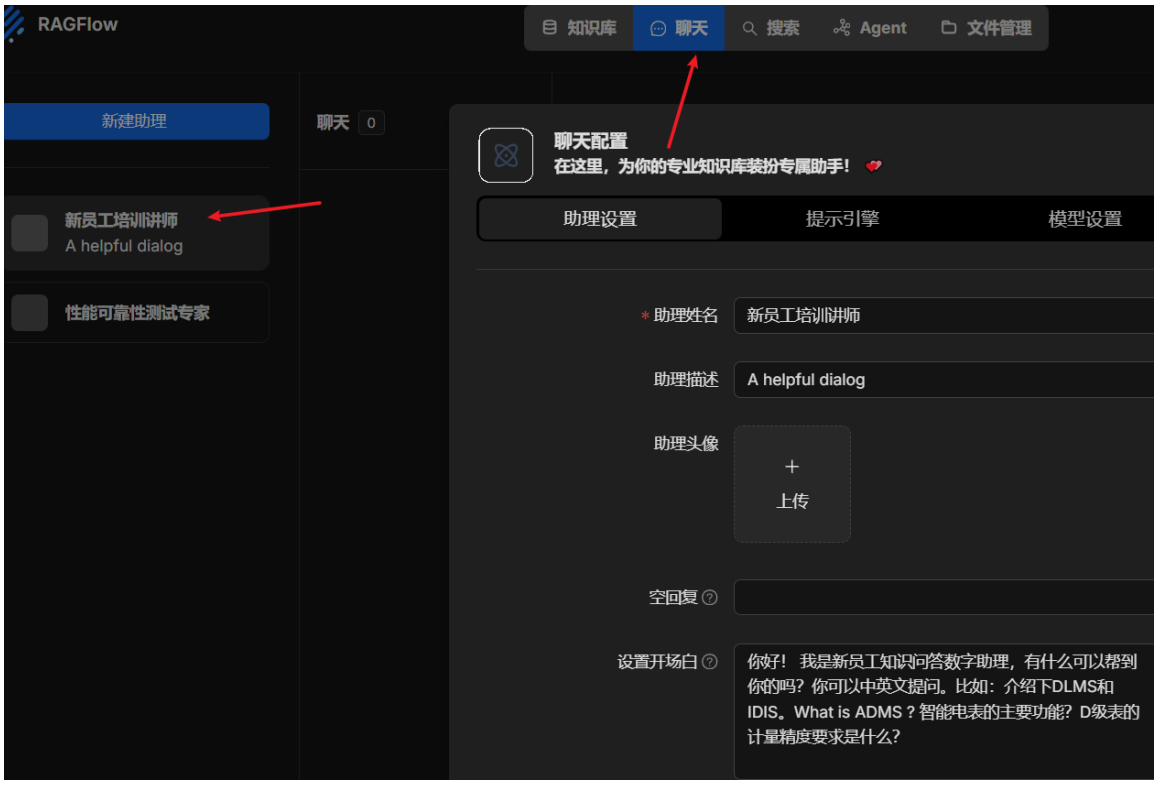

Figure 8: Create application

step 5: Publish app, and test, also can embedded the link to your own web.

```
<iframe
src="http://172.16.33.244:8080/chat/share?shared_id=da54fcfe
f9a311efa8a70242ac120006&from=chat&auth=
FiYWQ3MzMwZjgxMzExZWZhNmZlMDI0Mm"
style="width: 100%; height: 100%; min-height: 600px"
frameborder="0"
>
</iframe>
```

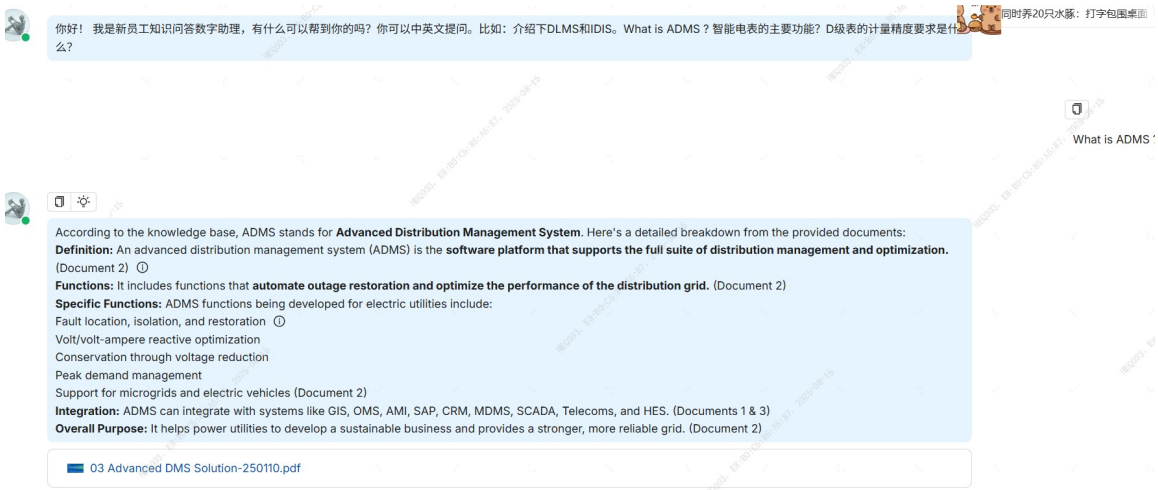

Figure 9: Test the Digital Human

2.4 Dify - low-code AI development framework

Run all the container images with Docker. Once they start normally, the service will be exposed on the default port 80 . During the first initialization you must set up a system-administrator account and password. Currently, the admin user name is : **q30china@gmail.com**, password: **hexing@2025**

```
# start the service
cd /data/git/dify
docker compose up -d

# stop the service
docker compose down
```

A modern, open-source,low-code AI development framework that lets you create, host, and scale AI agents (chatbots, virtual assistants, knowledge-base bots, workflow automators, etc.) without writing boilerplate ML code.

Figure 10: Dify login

Select models, make workflow and programming callback function of workflow

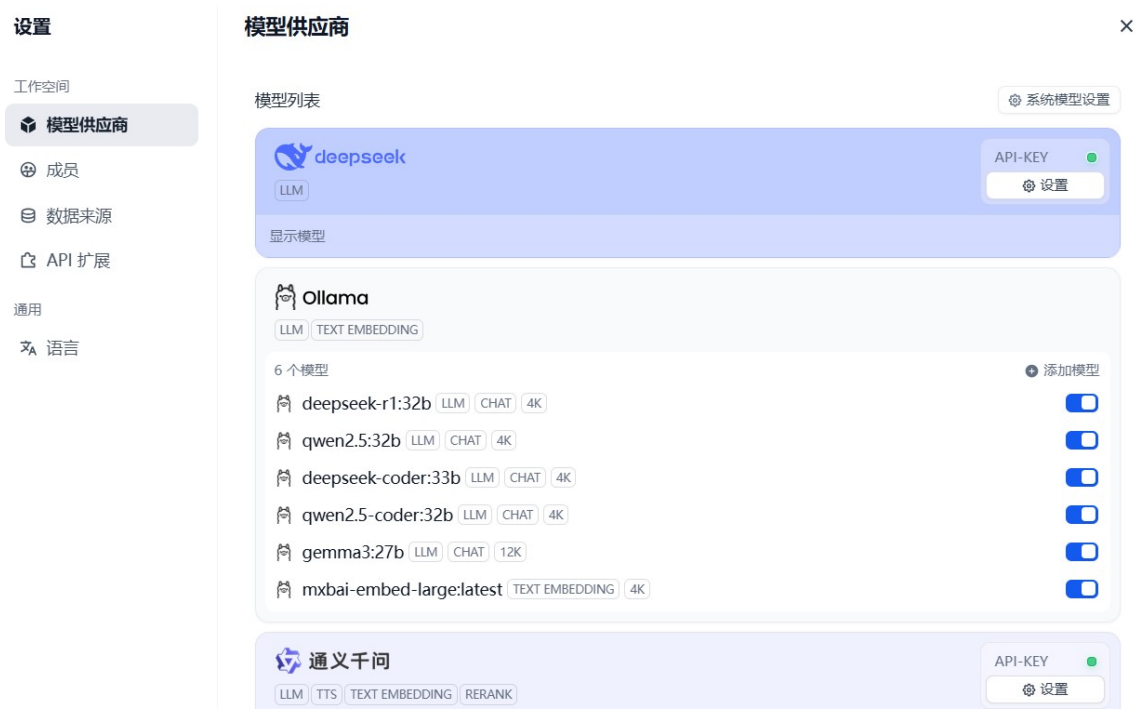

Figure 11: Configure models locally

```python
def call_dify_workflow(url, prompt, api_key, base_url):
    """
    调用dify中的工作流API
    :url: 网址字符串
    :prompt: 提示词字符串
    :api_key: 工作流API_KEY
    :base_url: API的调用地址
    """
    api_url = f"{base_url}/workflows/run"
    headers = {
        "Authorization": f"Bearer {api_key}",
        "Content-Type": "application/json"
    }
    payload = {
        "inputs": {
            "input_url": url,
            "prompt": prompt
        },
        "response_mode": "blocking",
        "user": "user-123"
    }
    records = []
    response = requests.post(api_url, headers=headers, data=json.dumps(payload))
    if response.status_code == 200:
        result = response.json()
        #print(json.dumps(result, ensure_ascii=False))
        records = process_json_data(result)
        df = pd.DataFrame(records)
        for _, row in df.iterrows():
            for i, column in enumerate(df.columns):
                value = row[i]
                print(f"{column}: {value}")

            print("\n")  # 每条记录之间加一个空行

        return records

    else:
        print(f"API调用失败，状态码：{response.status_code}")
        print(f"响应内容：{response.text}")
        return []
```

Figure 12: Call back function of workflow

example 1 : example of Auto acquisition of tenders from website agent

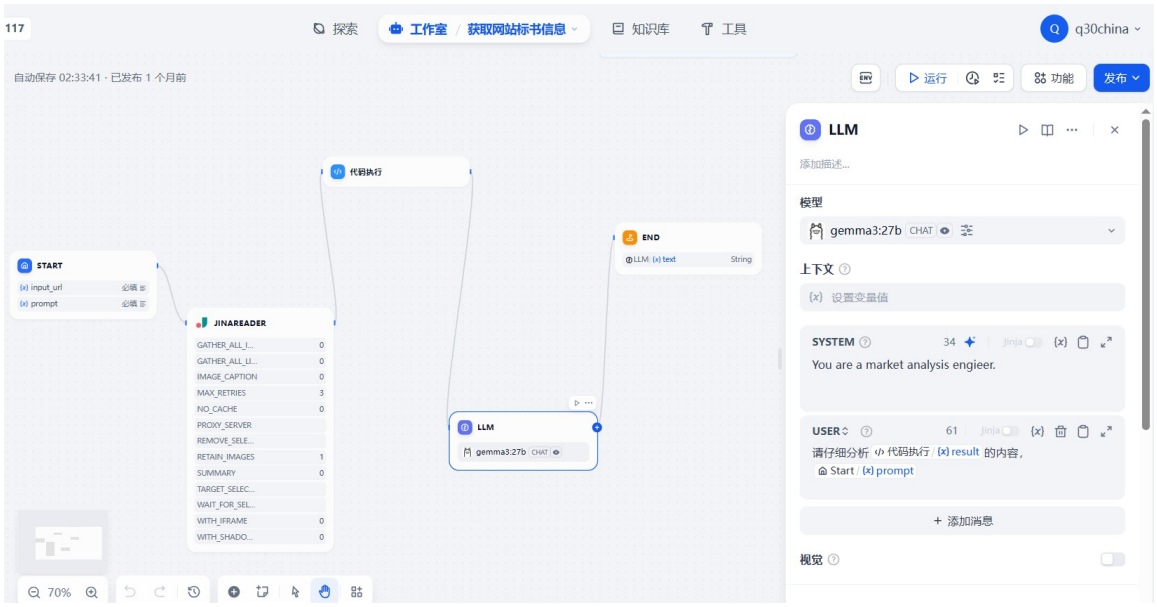

Figure 13: Workflow definition for tender acquisition

example 2 : example of An Testing Interview agent

Figure 14: Interview agent demo

2.5 PaddleOCR

PaddleOCR is an open-source, multi-language OCR toolkit developed on the PaddlePaddle deep learning framework.

Here we use the PaddleOCR toolkit to perform image recognition on the LCD display of an intelligent electric meter. The model employed is a self-trained model built upon the open-source baseline. By annotating 1,000 photos of the LCD, we train with the PaddleOCR toolkit to generate our own model. Then, through calls to the OCR-Server / OCR-Client software, we enable automated testing of the intelligent electric meter's LCD display.

```
# create and use python virtual enviroment
conda activate paddleocr

# be sure  paddlepaddle-gpu is installed
python -m pip install paddlepaddle-gpu==3.0.0b1 -i https://pypi.tuna.
   tsinghua.edu.cn/simple

# modify the config file
cd /data/git/PaddleOCR/configs/rec
vi ppocr.yml

# copy trainning data to direcory train (1239 pics)
cp * /data/git/PaddleOCR/train_data/Lcd/train/

# copy testing data to directory test (131 pics)
cp * /data/git/PaddleOCR/train_data/Lcd/test/

# start training
python tools/train.py -c configs/rec/ppocr.yml
```

Some parameters from ppocr.yml :

```
ems@aisrv: /data/git/PaddleOCR/configs/rec  ×  ems@aisrv: /data/g
Global:
  debug: false
  use_gpu: true
  epoch_num: 50
  log_smooth_window: 20
  print_batch_step: 10
  save_model_dir: ./output/v3_7-25
  save_epoch_step: 3
  eval_batch_step: [0, 10]
  cal_metric_during_train: true
  pretrained_model: ./output/v3_1-18/best_accuracy
  checkpoints:
  save_inference_dir:
  use_visualdl: false
  infer_img: doc/imgs_words/ch/word_1.jpg
  character_dict_path: ppocr/utils/en_dict.txt
  max_text_length: &max_text_length 25
  infer_mode: false
  use_space_char: true
  distributed: true
  save_res_path: ./output/rec/predicts_ppocrv3_en.txt


Optimizer:
  name: Adam
  beta1: 0.9
  beta2: 0.999
  lr:
    name: Cosine
    learning_rate: 0.001
    warmup_epoch: 5
  regularizer:
    name: L2
    factor: 3.0e-05


Architecture:
  model_type: rec
  algorithm: SVTR
  Transform:
  Backbone:
    name: MobileNetV1Enhance
    scale: 0.5
    last_conv_stride: [1, 2]
    last_pool_type: avg
  Head:
    name: MultiHead
    head_list:
      - CTCHead:
          Neck:
"ppocr.yml" [dos] 132L, 2799B
```

Figure 15: Part of ppocr.yml

Training process :

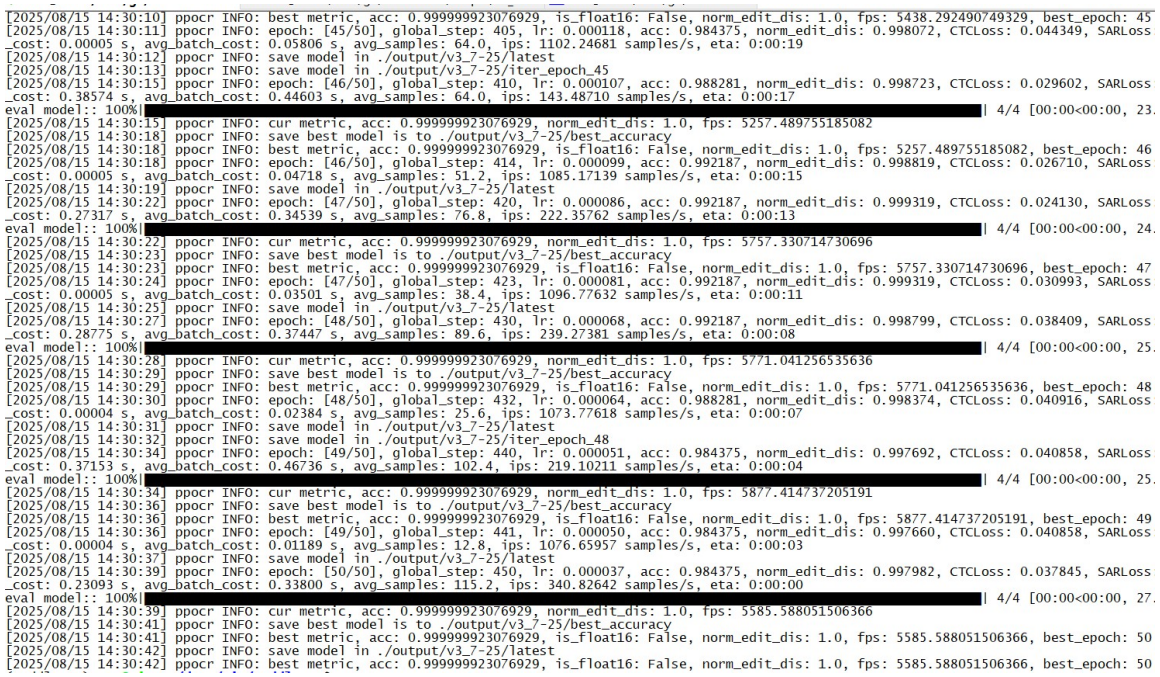

Figure 16: 50 epoch train process

2.6 AI4Test

AI4Test project main **objectives**: (2024.1.1-2024.12.1)

- Optimize Test Book Generation: Automatically Support the Test Engineering Team in designing the Test Books;
- Standardize Test Cases: Normalize Test Cases through a unified Common Test Case Language Model;
- Transform Unstructured Test Cases with AI: Utilize AI models to convert unstructured Test Cases into a standardized Common Language Model;
- AI-Driven Requirements-to-Test Cases Modeling: Create an AI model for predicting and automating test generation, prioritization, reporting test durations, and forecasting failure probabilities based on product requirements;
- New Test Cases Scenarios: Explore the generation of novel Test Cases from product requirements using large language models (LLMs).

How to start and use AI4Test ?

```
# create and use python virtual enviroment
conda activate ai4test

# Start AI4Test Webserver
cd /data/git/ai4test
python3 Infrastructure/app/main.py

# WebSite address
http://172.16.33.244:8882/
```

1/ The main user interface for AI4Test

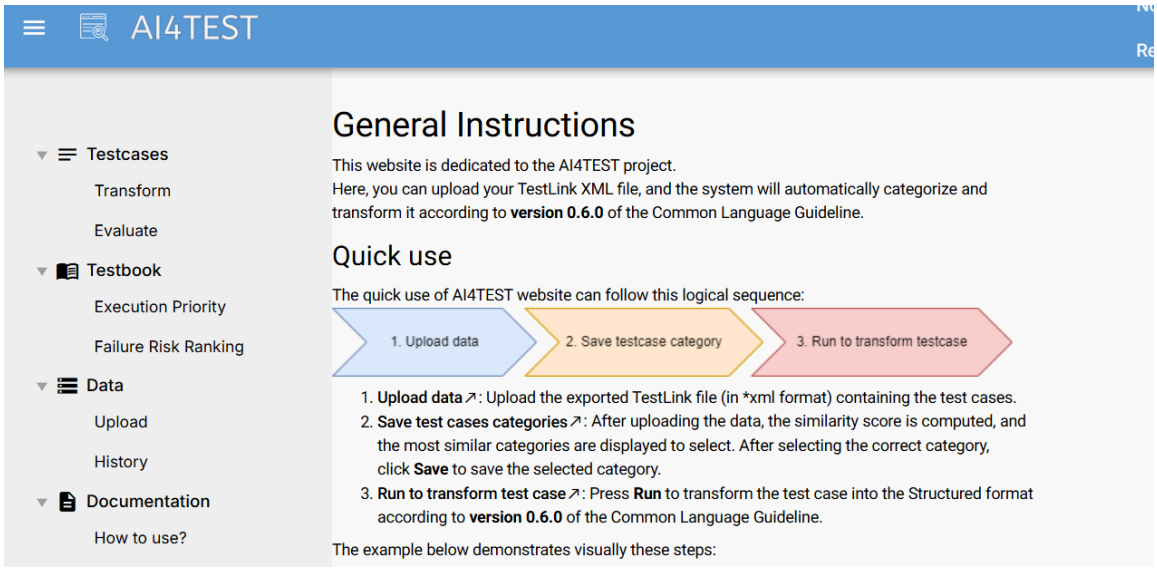

Figure 17: AI4Test main UI

/2 When using the tool for the first time, you need to upload test cases in XML format, which can be directly exported from TestLink. In addition, you can upload requirement documents—docx or xls files are supported.

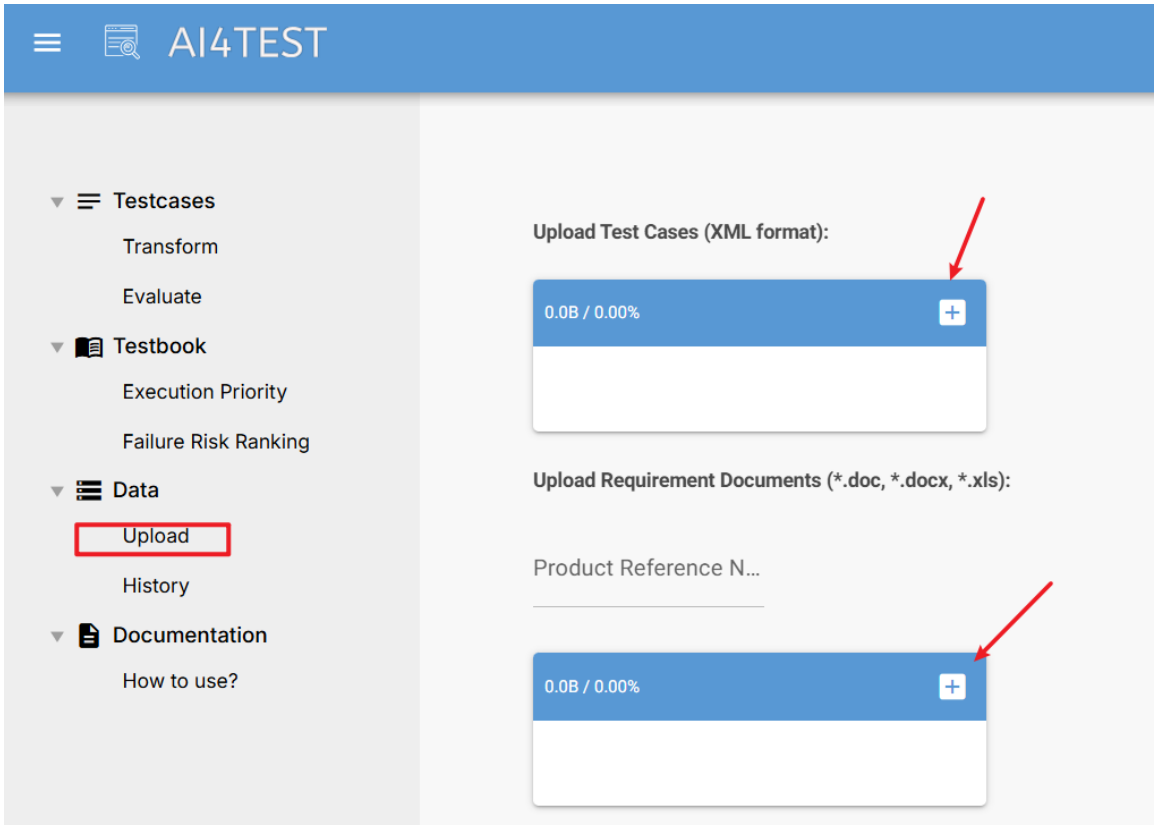

Figure 18: Upload files to AI4Test

3/ Loading requirements and test cases can be done by selecting them from the history of uploaded data

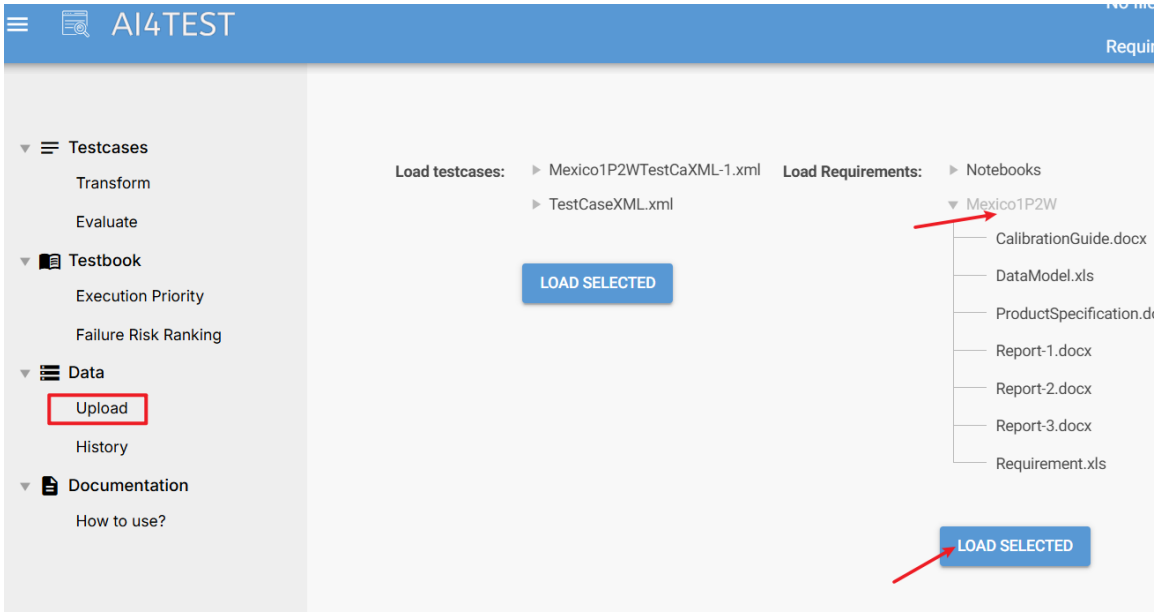

Figure 19: Loading requirements

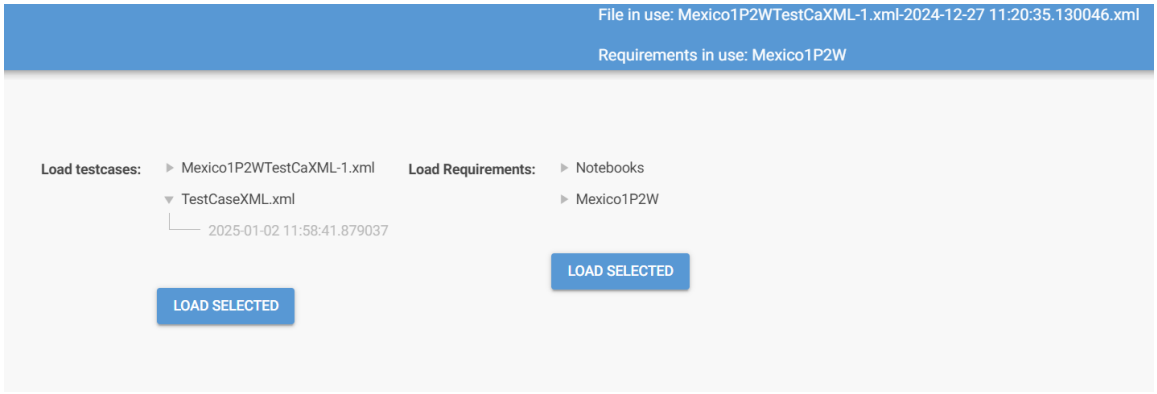

Figure 20: Loading test cases

4/ From the Transform section of the Testcases page in the left-hand tree, you can view the AI conversion results, whereby test cases exported from TestLink are transformed into structured test cases defined by the Common Test Language specification

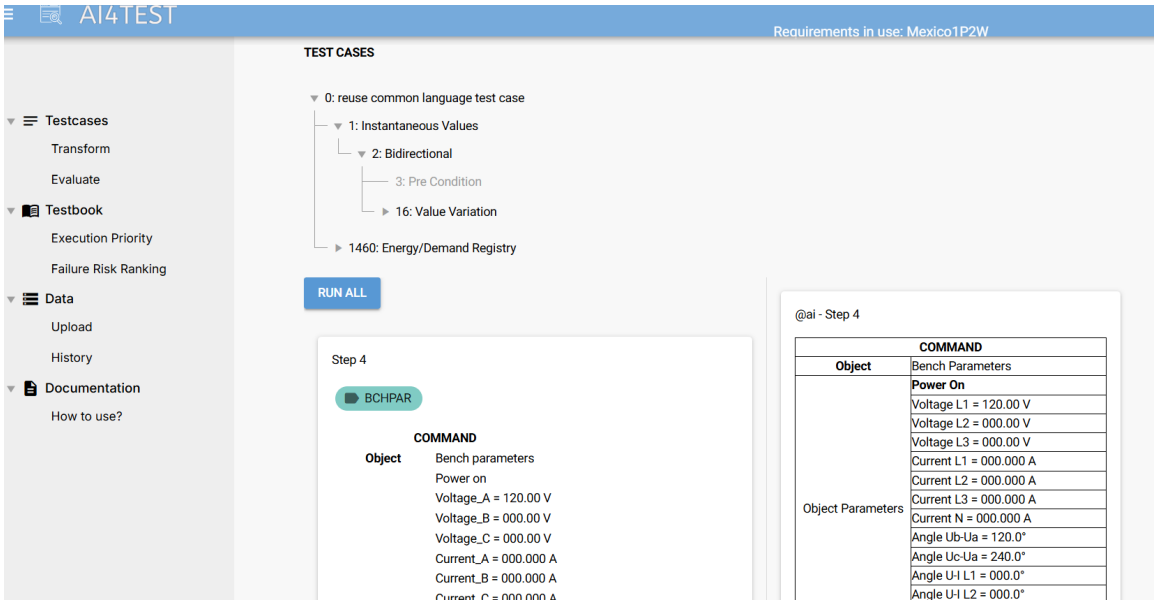

Figure 21: Test Case transformation

5/ If the AI-converted result is incorrect, you can manually correct it. You can select a different classification, and other categories are also available in the Additional section. Once you submit your feedback, the system will automatically retrain and learn from it.

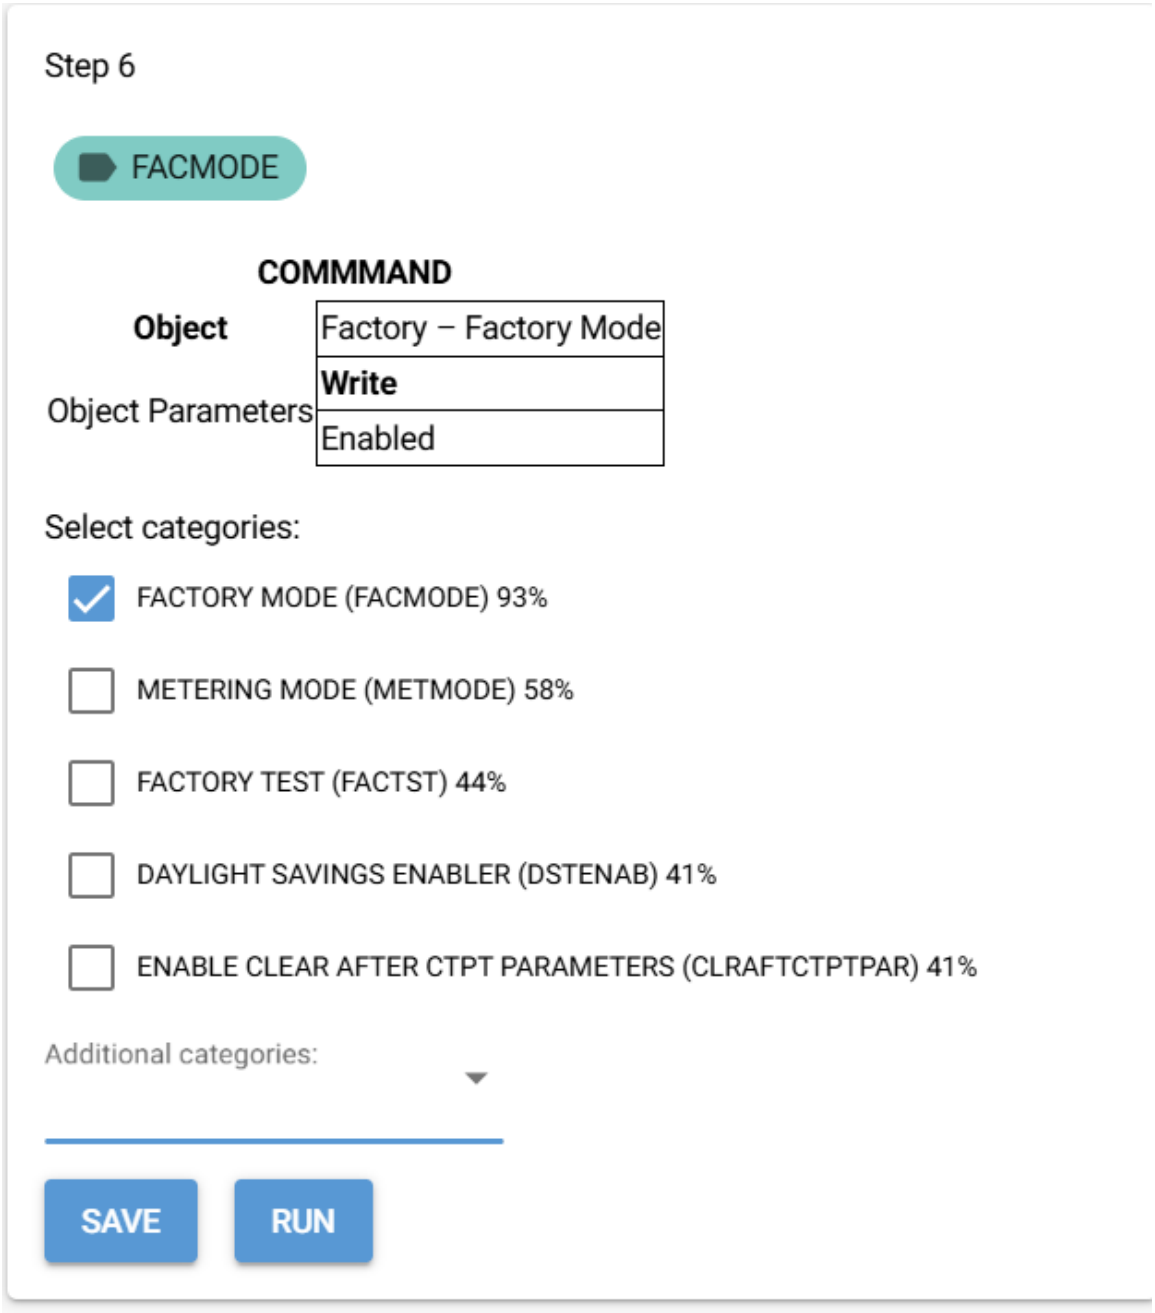

Figure 22: Manual adjustment of conversion results

6/ Each category is accompanied by a probability—the higher the value, the stronger the AI's recommendation.

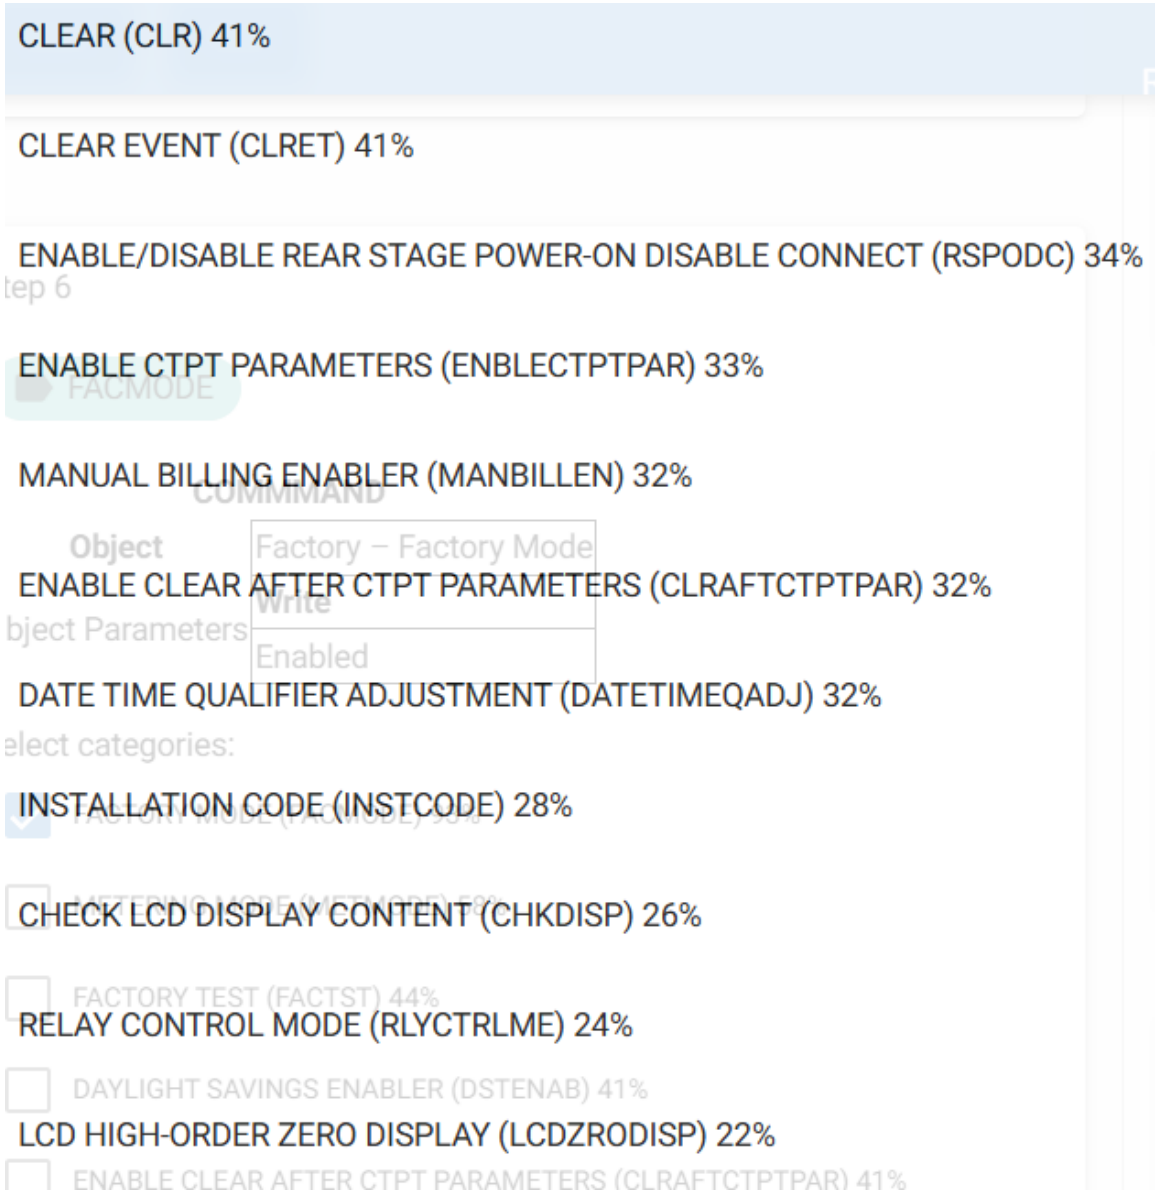

Figure 23: Execution step classification

7/ Sort the test cases by execution success rate based on their historical execution results.

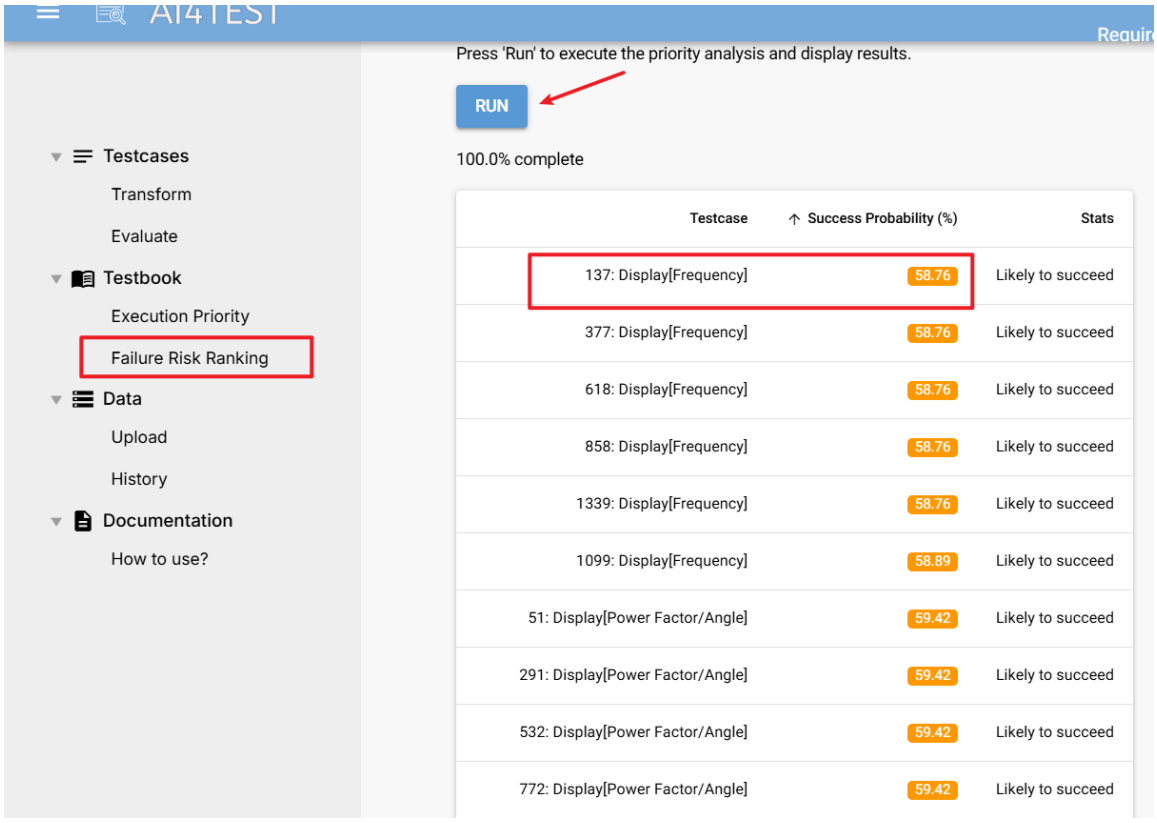

Figure 24: Test case execution success rate

3 Closing remarks

All of the contents in this document represnets my exploration in AI while learning, supplemented by references to various project documents. There may be some errors and omissions, so please feel free to contact me if you find any. I would also like to thank the members of the AI4Test team — **Francisco Alexandre** and **Wang Wen**—as well as **Gabriel** from the OCR project. Though I was involved in both projects from start to finish, I appreciate their contributions.